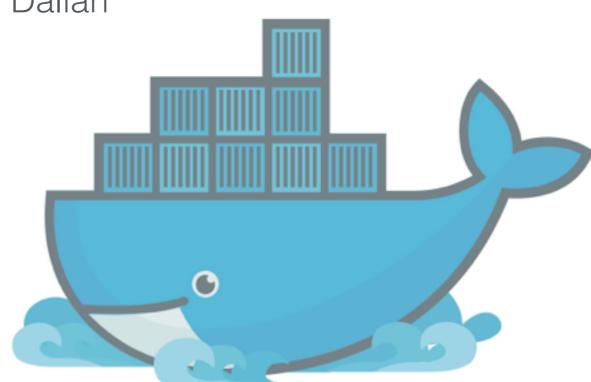
基于Docker的云处理服务架构实践



叶靖 (yejingx) @ 又拍云 杭州

2016.7.9 SF D-Day Dalian



6 UPJUN





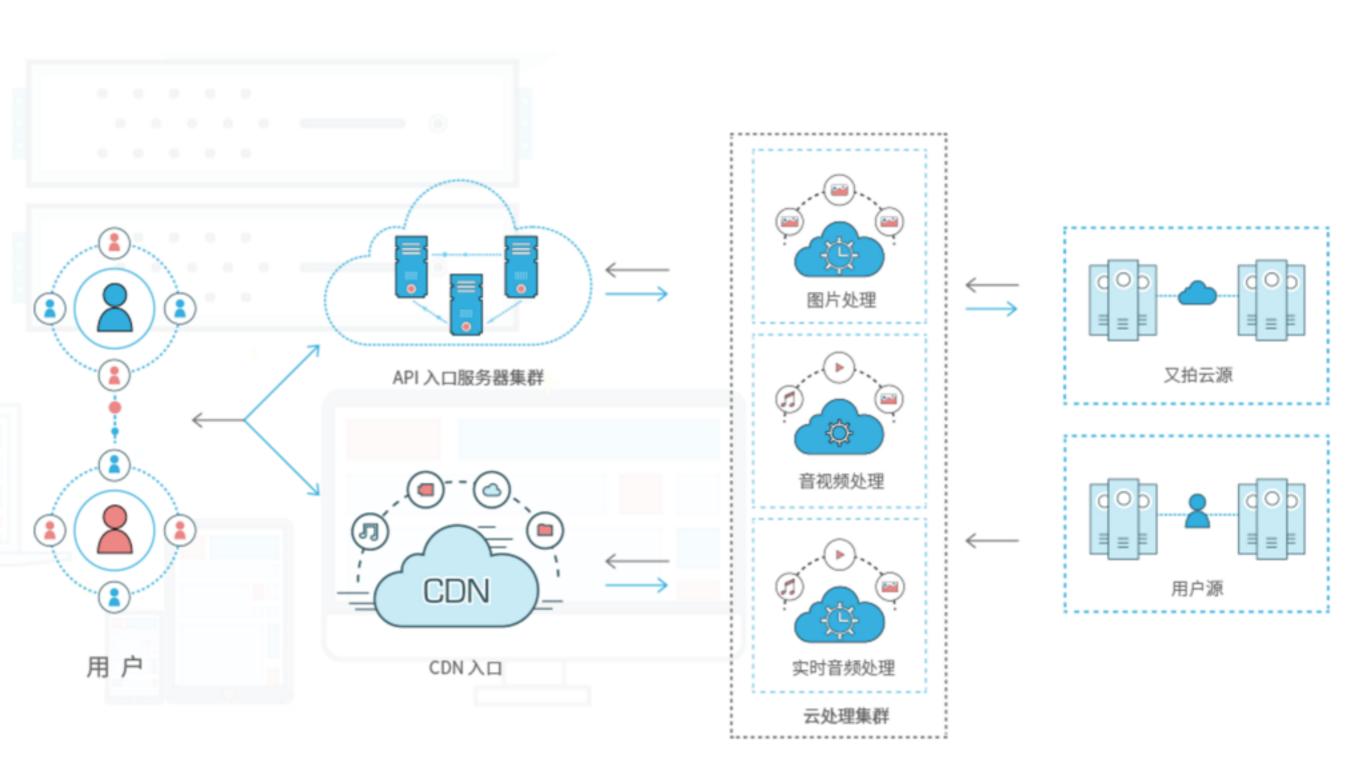






云处理在又拍





裸机部署



部署、升级复杂

环境乱

资源抢占

管理分散

资源利用率低

Docker 的优势



资源、环境隔离

部署、升级方便

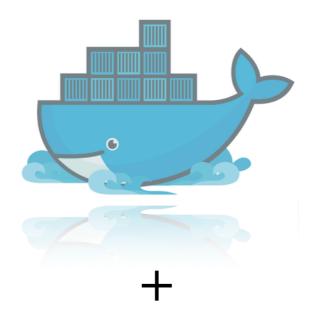
扩容方便

统一管理

社区活跃

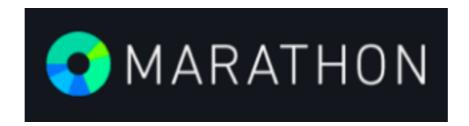
性能好





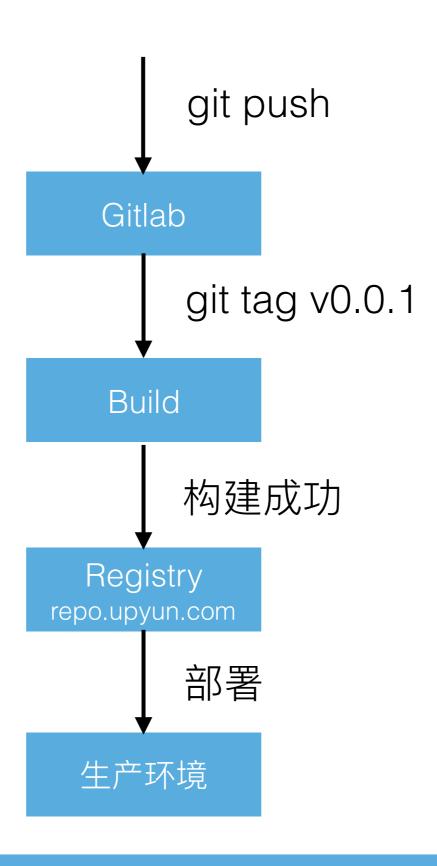






基于 Docker 的 CI





基于 Docker 的 CI



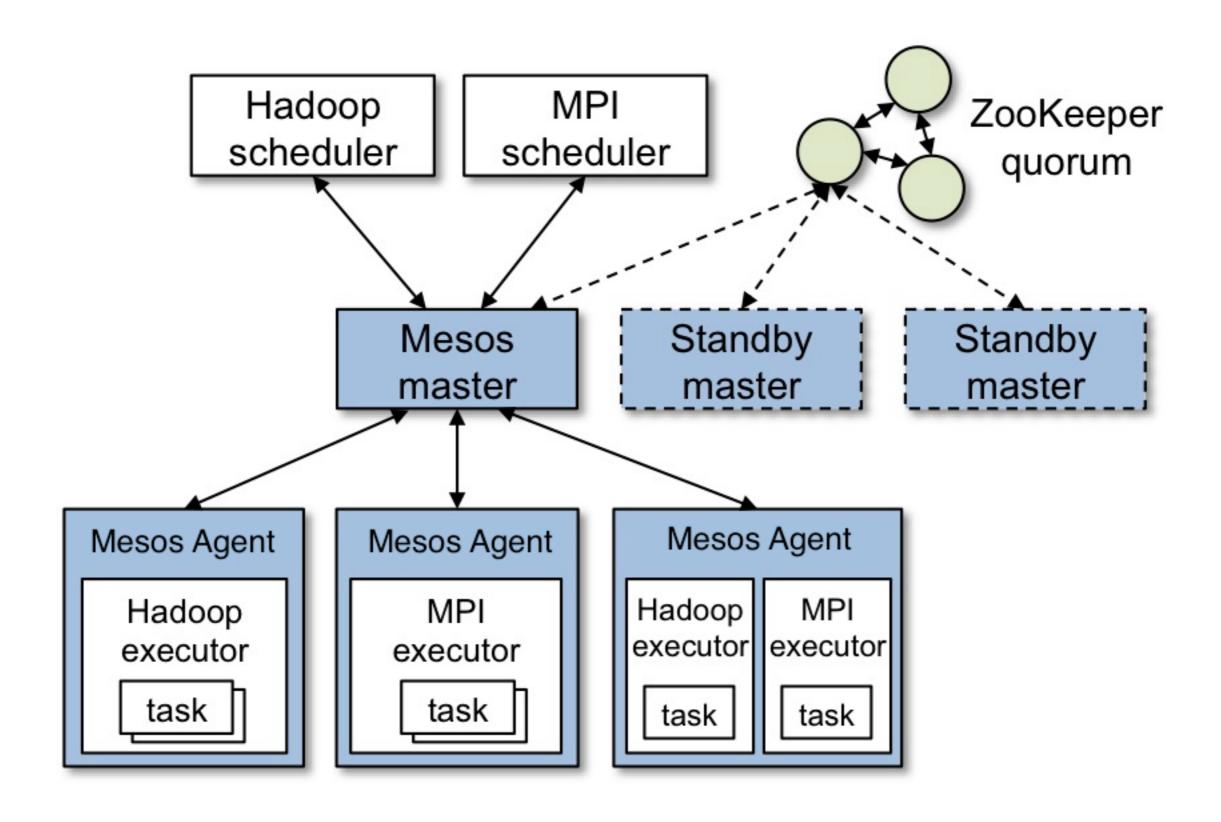
#.gitlab-ci.yaml

```
build:
    stage: build
    script:
        - img=repo.upyun.com:5043/alertman:$CI_BUILD_TAG
        - docker build -t $img .
        - docker push $img
        - docker rmi $img
        tags:
              - on-line-docker-builder
    only:
              - tags@consumers/alertman
```



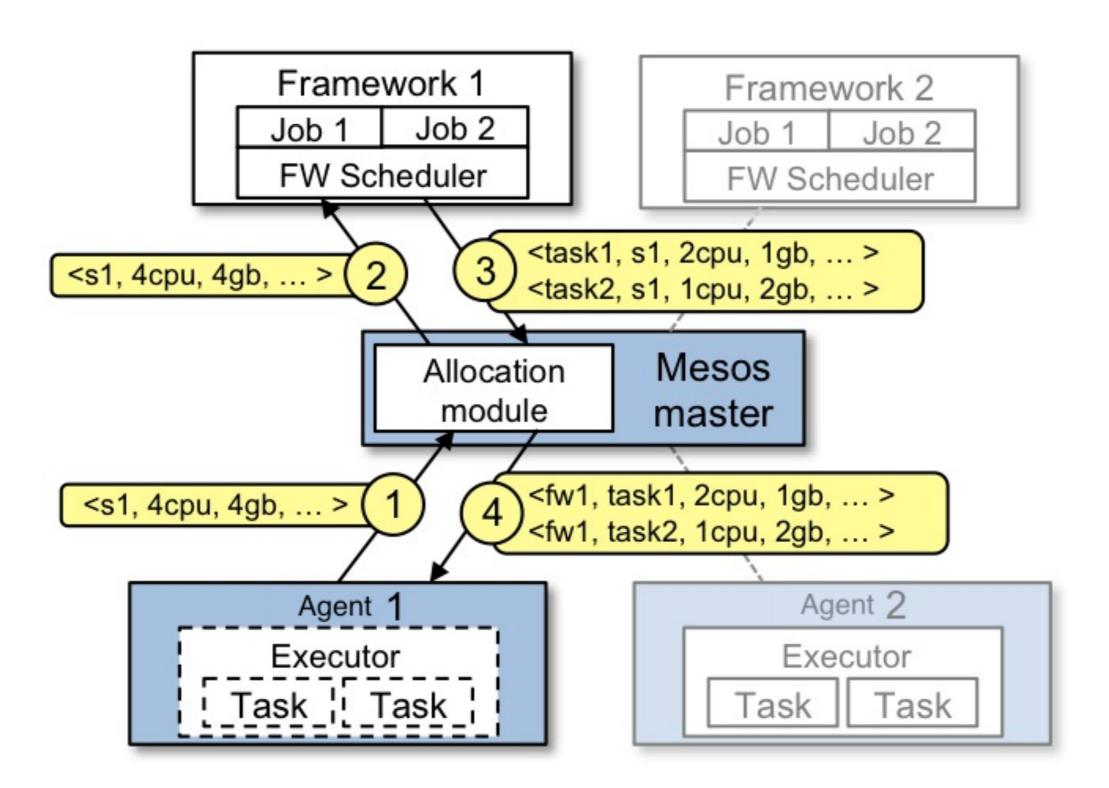
Mesos





Mesos





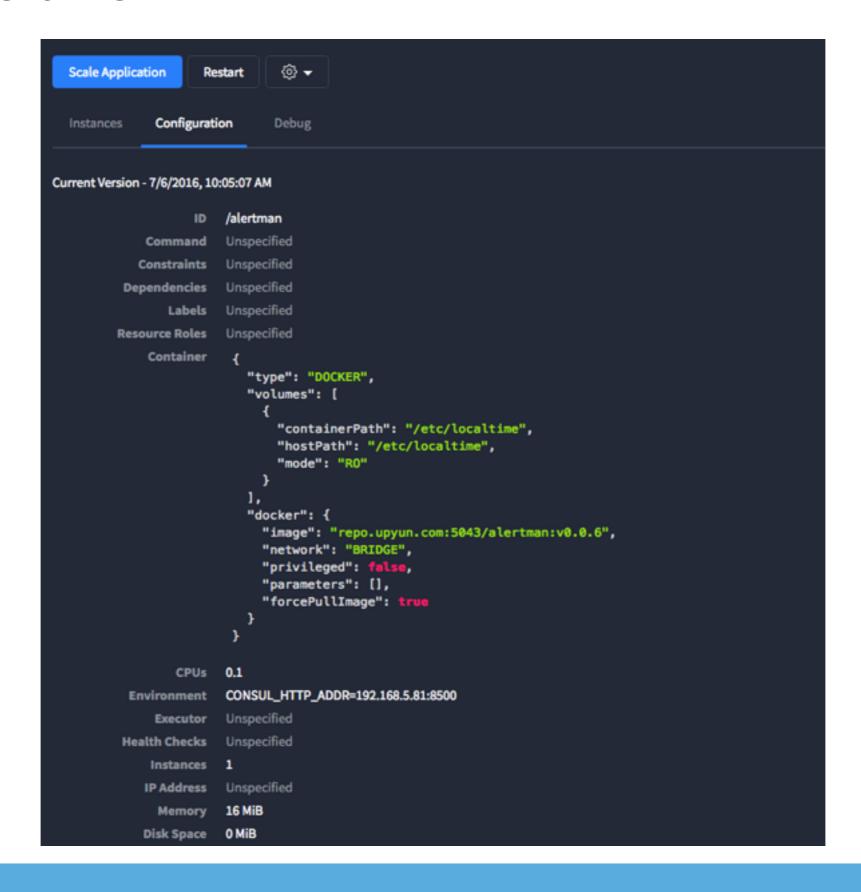
Marathon



MARATHON	Applications Deployments				About	API Reference	Documentation
Create	Applications				Se	earch all application	ns Q
STATUS 19	Name -	СРИ	Memory	Status	Running I	nstances Health	
☐ Deploying		0.1	16 MiB	Running		1 of 1	
Suspended 11 Delayed	api-aduit	0.0	0 B	Suspended		0 of 0	
Waiting	api-notify	1.0	256 MiB	⊘ Running		1 of 1	
HEALTH Healthy	api-notify-host	0.1	256 MiB	⊗ Running		1 of 1	
Unhealthy	apiaduit	0.0	0 B	Suspended		0 of 0	
Unknown 19		0.1	16 MiB	⊗ Running		1 of 1	
LABEL ▼	⇔ -	0.1	16 MiB	⊗ Running		1 of 1	
		0.2	128 MiB	⊘ Running		2 of 2	
	⇔	0.1	16 MiB	⊗ Running		1 of 1	
	docker-es-agent	0.5	80 MiB	⊘ Running		5 of 5	
	gitlab-runner	0.0	0 B	Suspended		0 of 0	

Marathon













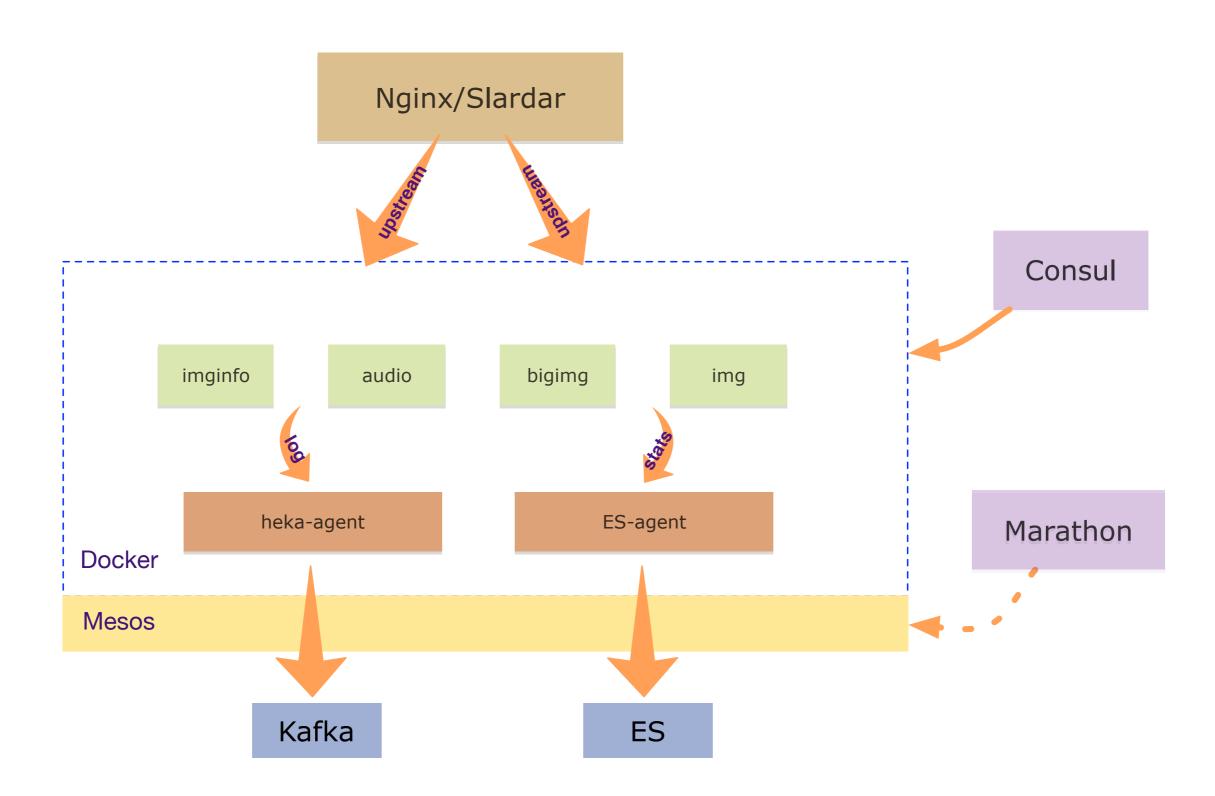
光有这些是 不够的!

我们还需要



- ▶服务路由
- ▶配置管理
- ▶日志收集
- ▶ 服务状态
- ▶监控报警
- **....**





动态服务路由



HAProxy?

动态服务路由



上游列表如何改变?

新服务如何添加?









lua-resty-logger-socket

lua-resty-httpipe lua-resty-checkups

lua-resty-limit-traffic

lua-resty-httproxy



以 Host 区分服务:

```
$ curl -T cat.jpg 192.168.1.155:3130 \
    -H "Host: imageinfo"
$ curl -T cat.jpg 192.168.1.155:3130 \
    -H "Host: imgprocess"
```



通过 HTTP 接口动态更新 upstream 列表:



```
▼ "cls:node-dev.upyun.com": [
            "server": "node-dev.upyun.com:10.0.5.108:4001",
            "msg": null,
            "status": "ok",
            "lastmodified": "2016-07-05 16:23:48",
            "fail num": 0
            "server": "node-dev.upyun.com:10.0.5.109:4001",
            "msg": "connection refused",
            "status": "err",
            "lastmodified": "2016-07-06 14:50:22",
            "fail num": 1
```

Consul 配置管理





Service Discovery



Failure Detection



Multi Datacenter



Key/Value Storage

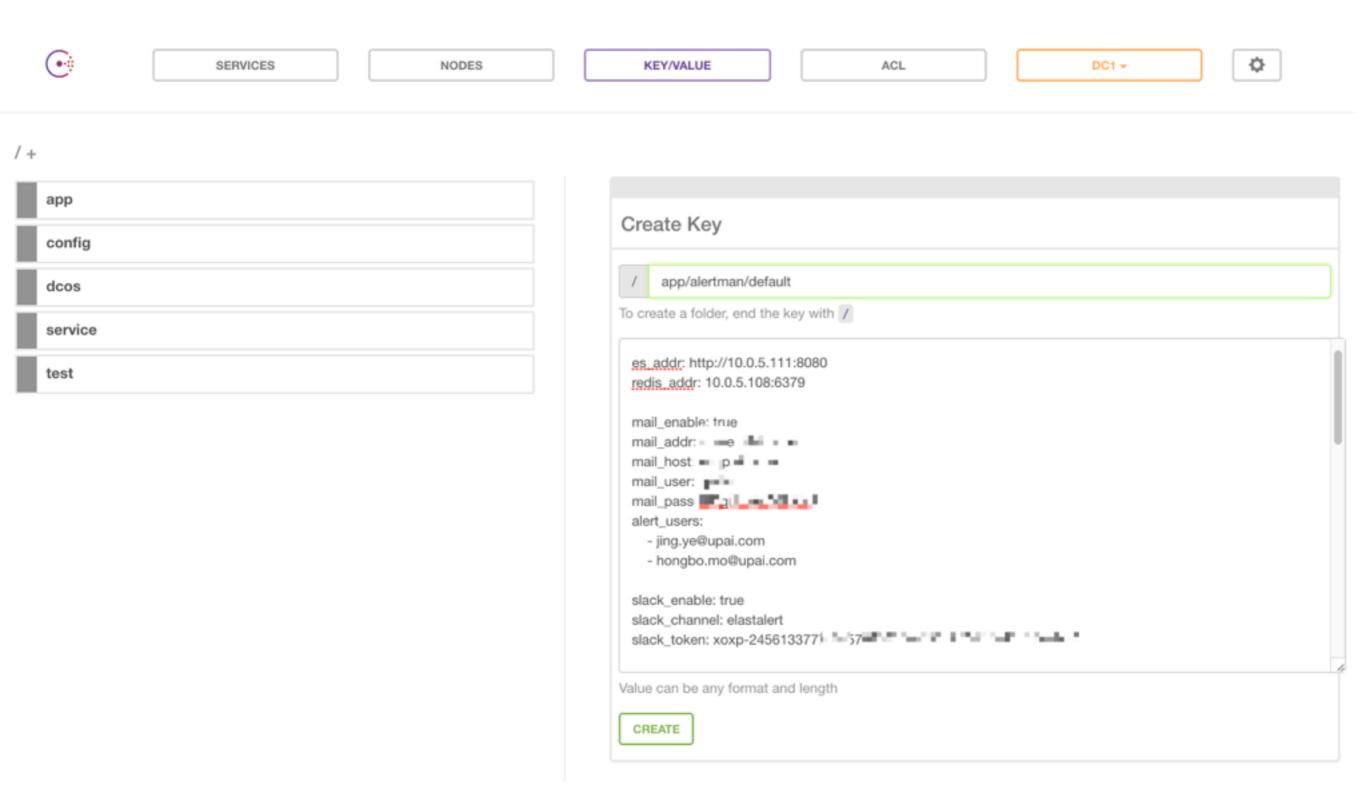
Consul 配置管理



```
Terminal
admin@hashicorp: curl -X PUT -d 'bar' http://localhost:8500/v1/kv/foo
true
admin@hashicorp: curl http://localhost:8500/v1/kv/foo
       "CreateIndex": 100,
       "ModifyIndex": 200,
       "Key": "foo",
       "Flags": 0,
       "Value": "YmFy"
admin@hashicorp: echo "YmFy" | base64 --decode
bar
admin@hashicorp:
```

Consul 配置管理





日志收集



日志量大、频率高



分布在多台机器



经常迁移



按服务分割

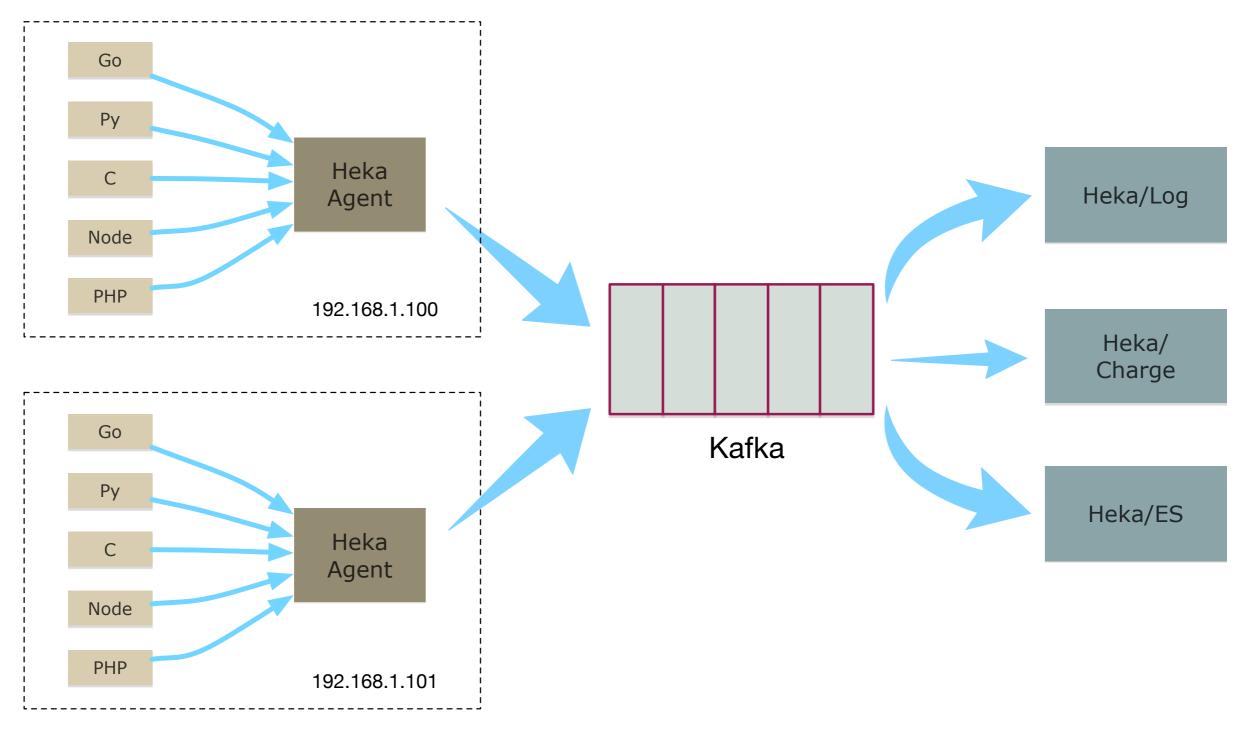


多种用途

日志收集



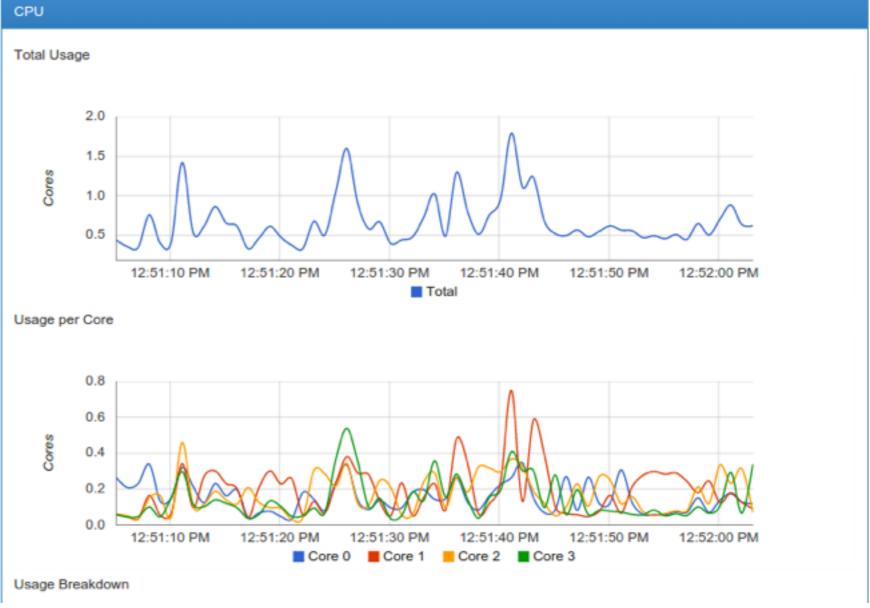
Heka-Agent + Heka Plugin

















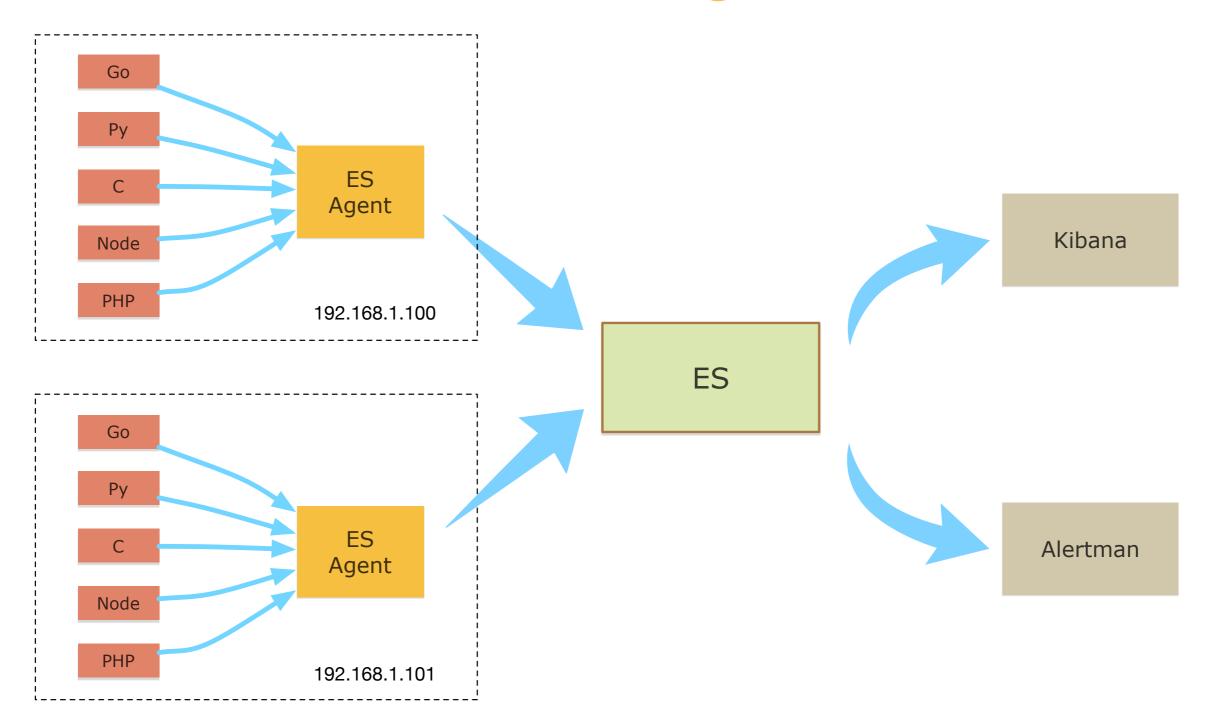


0.8 0.6 0.2 0.0 12:51:10 PM 12:51:20 PM 12:51:30 PM 12:51:40 PM 12:51:50 PM 12:52:00 PM

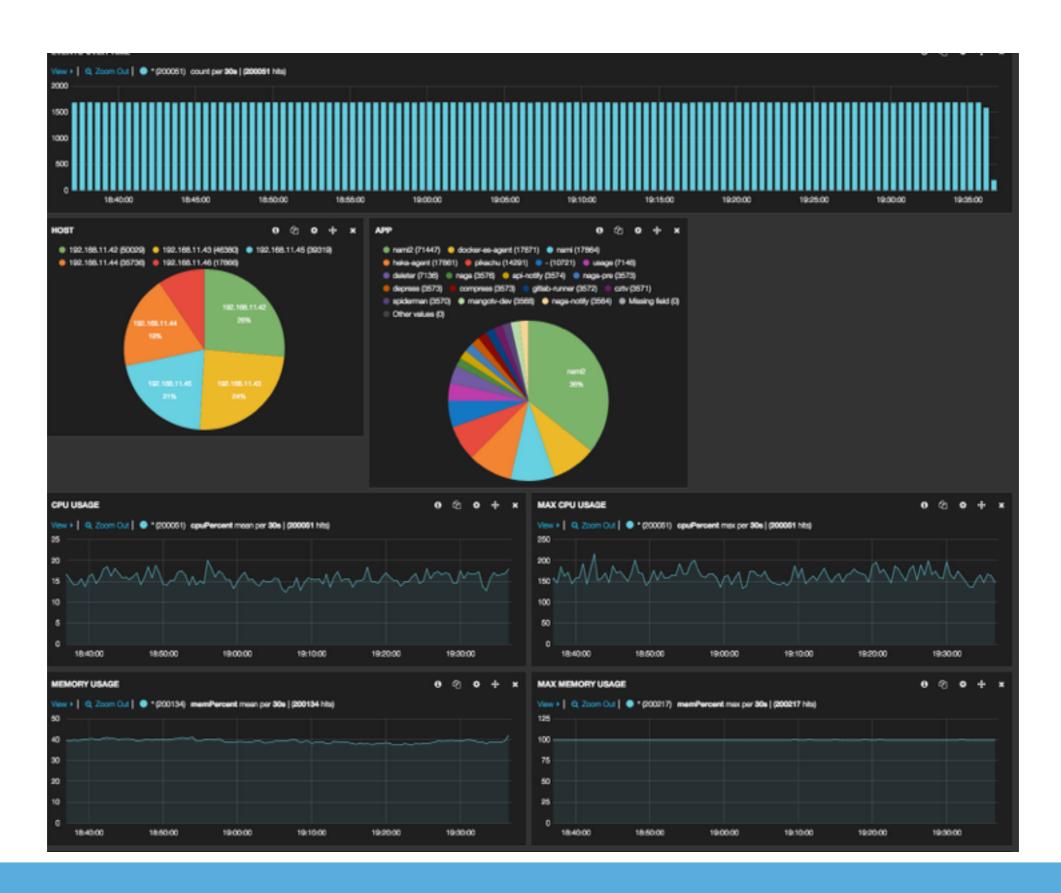
Core 0 Core 1 Core 2 Core 3



docker-es-agent







监控报警



Nagios°

VS

Alertman



监控报警



- ▶ 不需要大而全的监控
- ▶ 异常点分析
- ▶自动故障恢复
- ▶ 业务相关性大

监控报警 - 自己造轮子

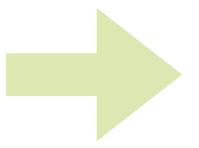


Nginx 日志

ES 数据

队列

数据分析



Alertman

监控报警 - 实例



```
2016-06-02 14:09:27,461 - status - DEBUG - got hits of status 504: [(u'192.168.13.245:6501', 438), (u'192.168.13.243:6501', 24), (u'192.168.13.34:6501', 23), (u'192.168.13
.242:6501', 15), (u'192.168.13.36:6501', 15), (u'192.168.13.105:6501', 14), (u'192.168.11.19:6501', 14), (u'192.168.11.18:6501', 14), (u'192.168.11.17:6501', 14), (u'192.168.11.17), (u'192.168.11.17), (u'192.168.11.17), (u'192.168.11.17), (u'192.16
.13.237:6501', 12), (u'192.168.11.22:6501', 12), (u'192.168.11.21:6501', 12), (u'192.168.13.107:6501', 12), (u'192.168.13.106:6501', 12), (u'192.168.11.27:6501', 11), (u'1
.168.5.205:6501', 7), (u'192.168.5.208:6501', 7), (u'192.168.5.200:6501', 6), (u'192.168.5.209:6501', 5), (u'192.168.5.201:6501', 3), (u'192.168.5.208:6501 : 192.168.13.24
2016-06-02 14:09:27,464 - status - DEBUG - stddev_from_average mean: 23.0526315789, stddev: 69.348297342
2016-06-02 14:09:27,465 - status - INFO - anomalous nodes of status 504: [(u'192.168.13.245:6501', 438)]
2016-06-02 14:09:27,467 - status - INFO - try removing [(u'192.168.13.245:6501', 438)] from slardar_test
2016-06-02 14:09:42,386 - usage - INFO - nami usage is 0.00
2016-06-02 14:09:42,433 - usage - INFO - nami2 usage is 45.41
2016-06-02 14:09:42,616 - status - DEBUG - 502 count is 95(<100), sleep 15s and try later
2016-06-02 14:09:42,749 - status - DEBUG - 504 count is 240(>80), checking hists ...
2016-06-02 14:09:48,638 - status - DEBUG - got hits of status 504: [(u'192.168.13.245:6501', 451) (u'192.168.13.243:6501', 25) (u'192.168.13.244:6501', 22), (u'192.168.1
.20:6501', 15), (u'192.168.13.105:6501', 15), (u'192.168.13.242:6501', 15), (u'192.168.11.16:6501', 15), (u'192.168.11.25:6501', 14), (u'192.168.11.19:6501', 14), (u'192.168.11.25:6501', 15)
.13.101:6501', 13), (u'192.168.13.106:6501', 13), (u'192.168.13.237:6501', 12), (u'192.168.11.27:6501', 12), (u'192.168.13.107:6501', 12), (u'192.168.11.30:6501', 12), (u'
2.168.11.28:6501', 8), (u'192.168.5.205:6501', 7), (u'192.168.5.208:6501', 6), (u'192.168.5.200:6501', 5), (u'192.168.5.201:6501', 4), (u'192.168.5.208:6501 : 192.168.13.2
2016-06-02 14:09:48,641 - status - DEBUG - c+ddou from guargaa maar: 22 7260421052 - c+ddou: 71 4016746755
2016-06-02 14:09:48,642 - status - INFO anomalous nodes of status 504: [(u'192.168.13.245:6501', 451)]
2016-06-02 14:09:48,643 - status - INFO try removing [(u'192.168.13.245:6501', 451)] from slardar_test
2016-06-02 14:09:57,589 - usage - INFO - numt usage is
2016-06-02 14:09:57,714 - usage - INFO - nami2 usage is 46.03
2016-06-02 14:10:03,796 - status - DEBUG - 502 count is 94(<100), sleep 15s and try later
2016-06-02 14:10:03,886 - status - DEBUG - 504 count is 239(>80), checking hists ...
2016-06-02 14:10:09,323 - status - DEBUG - got hits of status 504: [(u'192.168.13.245:6501', 447), (u'192.168.13.243:6501', 25), (u'192.168.13.244:6501', 22), (u'192.168.1
.20:6501', 15), (u'192.168.13.105:6501', 15), (u'192.168.13.242:6501', 15), (u'192.168.11.25:6501', 14), (u'192.168.11.19:6501', 14), (u'192.168.11.16:6501', 14), (u'192.168.11.19:6501', 15), (u'192.168.11.19:6501', 14), (u'192.168.11.19:6501', 14), (u'192.168.11.19:6501', 14), (u'192.168.11.19:6501', 14), (u'192.168.11.19:6501', 14), (u'192.168.11.19:6501', 15), (u'192.168.11.19:6501', 15), (u'192.168.11.19:6501', 14), (u'192.168.11.19:6501', 15), (u'192.168.11.19:6501', 15), (u'192.168.11.19:6501', 14), (u'192.168.11.19:6501', 15), (u'192.168.11.19:6501', 15), (u'192.168.11.19:6501', 15), (u'192.168.11.19:6501', 15), (u'192.168.11.19:6501', 15), (u'192.168.11.19:6501', 15), (u'192.168.11.19:6501', 14), (u'192.168.11.19:6501', 15), (u'192.168.11.19), (u'192.168.11.19), (u'192.168.11.19), (u'192.168.11.19), (u'192.168.11.19), (u'192.168.11.19), (u'192.168.11.19), (u'192.168.11.19), (u'192.168.11.19), (
.13.101:6501', 13), (u'192.168.13.106:6501', 13), (u'192.168.13.237:6501', 12), (u'192.168.11.27:6501', 12), (u'192.168.13.107:6501', 12), (u'192.168.11.30:6501', 12), (u'
2.168.11.28:6501', 8), (u'192.168.5.205:6501', 7), (u'192.168.5.208:6501', 6), (u'192.168.5.200:6501', 5), (u'192.168.5.201:6501', 4), (u'192.168.5.208:6501 : 192.168.13.2
2016-06-02 14:10:09,327 - status - DEBUG - stddev_from_average mean: 23.5263157895, stddev: 70.7675054619
2016-06-02 14:10:09,328 - status - INFO - anomalous nodes of status 504: [(u'192.168.13.245:6501', 447)]
2016-06-02 14:10:09,329 - status - INFO - try removing [(u'192.168.13.245:6501', 447)] from slardar_test
2016-06-02 14:10:12,800 - usage - INFO - nami usage is 0.00
2016-06-02 14:10:12,905 - usage - INFO - nami2 usage is 46.88
2016-06-02 14:10:24,490 - status - DEBUG - 502 count is 96(<100), sleep 15s and try later
2016-06-02 14:10:24,550 - status - DEBUG - 504 count is 240(>80), checking hists ...
2016-06-02 14:10:28,154 - usage - INFO - nami usage is 0.00
2016-06-02 14:10:28,329 - usage - INFO - nami2 usage is 47.43
```

遇到的坑



kernel: nf_conntrack: table full, dropping packet

由于 Docker 使用了 NAT 网络,在高并发流量下,追踪连接状态的表满导致内核丢包。

解决:

net.netfilter.nf_conntrack_max = 655350

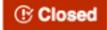
遇到的坑



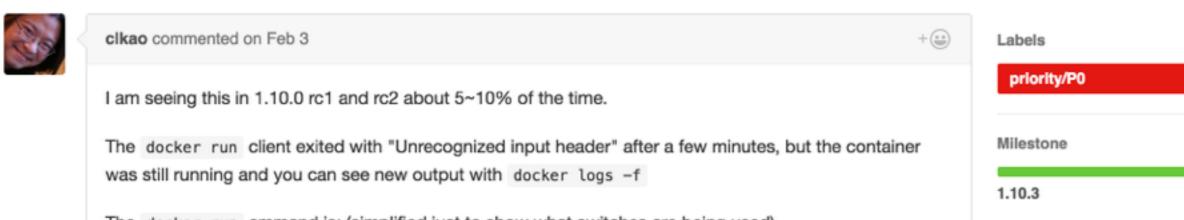
Mesos 里看到容器频繁重启,但 docker ps 看 旧进程却依然在

"Unrecognized input header" interrupting docker client #19950

New issue



(F) Closed clkao opened this issue on Feb 3 · 58 comments



The docker run cmmand is: (simplified just to show what switches are being used)

docker run --entrypoint /sbin/smell-baron --name somename -h somehost --volumes-from somevo

Milestone	Milestone 1.10.3		priority/P0	
		1.10.3	Milestone	
	1.10.3			

https://github.com/docker/docker/issues/19950

遇到的坑



.



欢迎加入又拍云

hr@upai.com

Q & A